

New York Times Article Abstract Analysis using Hadoop and NLTK

By: Lucas Nunno (lnunno@cs.unm.edu)

Part 1: Data Acquisition

Data acquisition

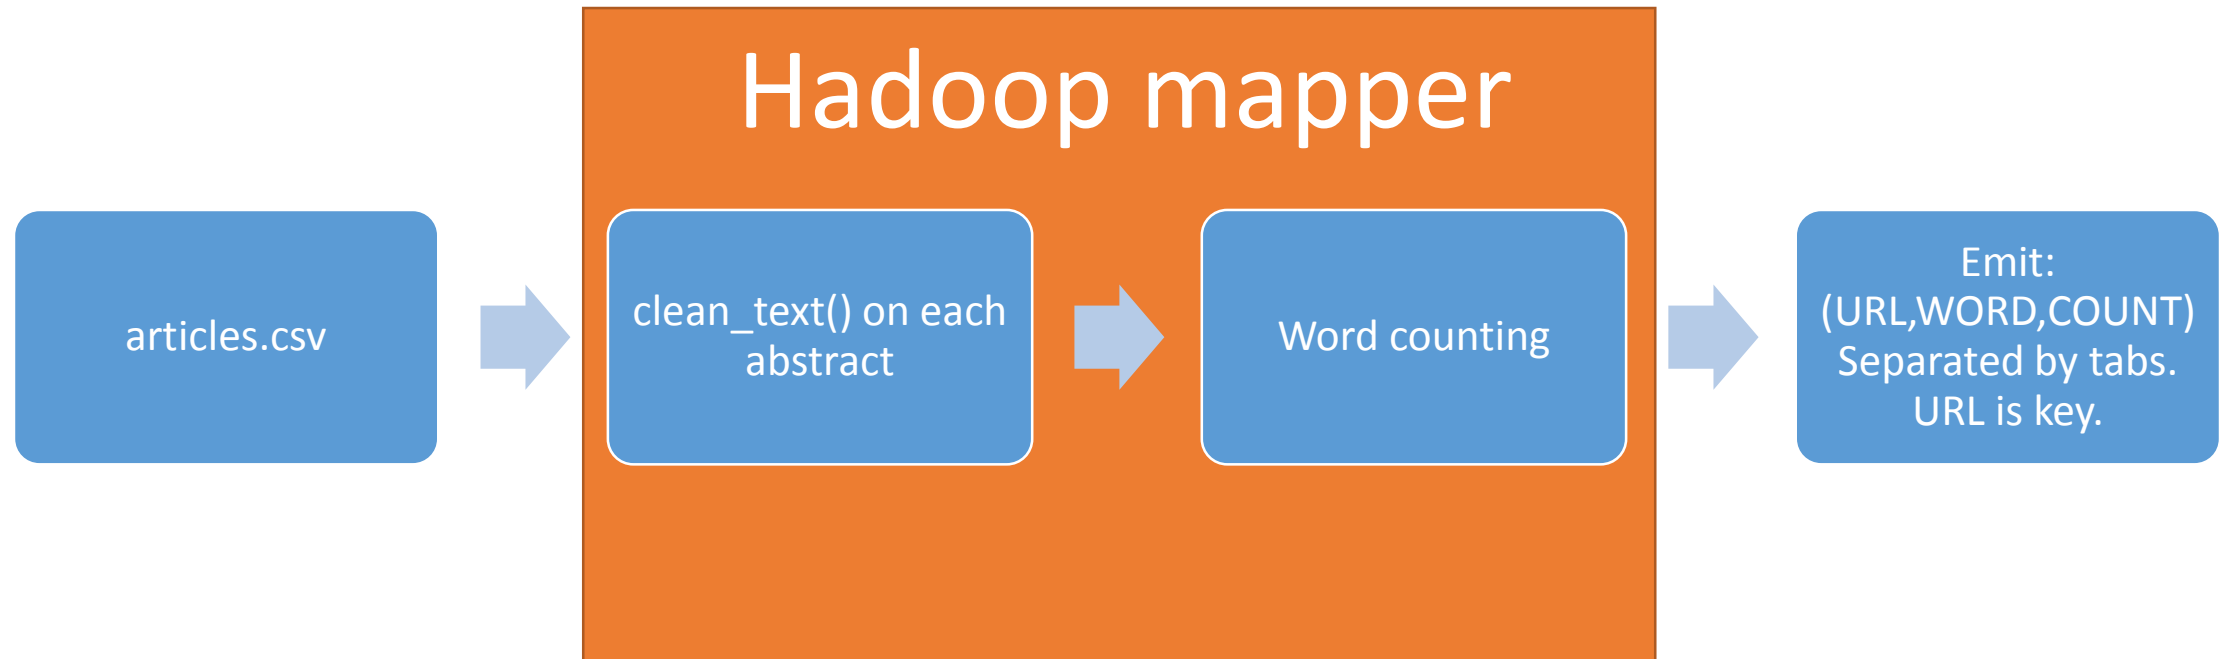
- Used the python **requests** module.
 - Used the offset parameter to load new pages of abstracts and slept 1/8th of a second between each request to abide by the NYT API terms of use.
- Loaded JSON response into python dictionary and then exported as a single large JSON file containing all the articles and all metadata.
(~40,000)
- In a separate script, I export this JSON data to a CSV file with the docIDs, URLs, and abstracts.
 - This is also where I check for **duplicates**. I have a set of URLs that the exporter has seen, if this URL is in this set the program prints a warning and does not export it.

Part 2: Preprocessing and tf-idf

Preprocessing

- Used the python natural language toolkit (NLTK) module for most of the preprocessing tasks. The algorithm is as follows:
 1. Convert text to lowercase.
 2. Remove punctuation and numbers.
 - Simple regex substitution:
`remove_pattern = re.compile(r'^a-z\s')`
 3. Remove stopwords.
 - See: `nltk.corpus.stopwords`
 4. Stem all the remaining words.
 1. See: <http://www.nltk.org/api/nltk.stem.html#module-nltk.stem.porter>
 5. Output the cleaned abstract.




How is this parallelized in Hadoop?



tf-idf

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$






- Dictionary is constructed of each document, the words it contains, and the frequencies of these words – this is all provided from the mapper.
- Frequency matrix is constructed for all documents.
 - Items without entries are filled with zeroes.
- Augmented (normalized) frequency matrix is calculated.
 - This removes the bias for longer documents.

$$tf(t, d) = 0.5 + \frac{0.5 \times f(t, d)}{\max\{f(w, d) : w \in d\}}$$

tf-idf (contd.)

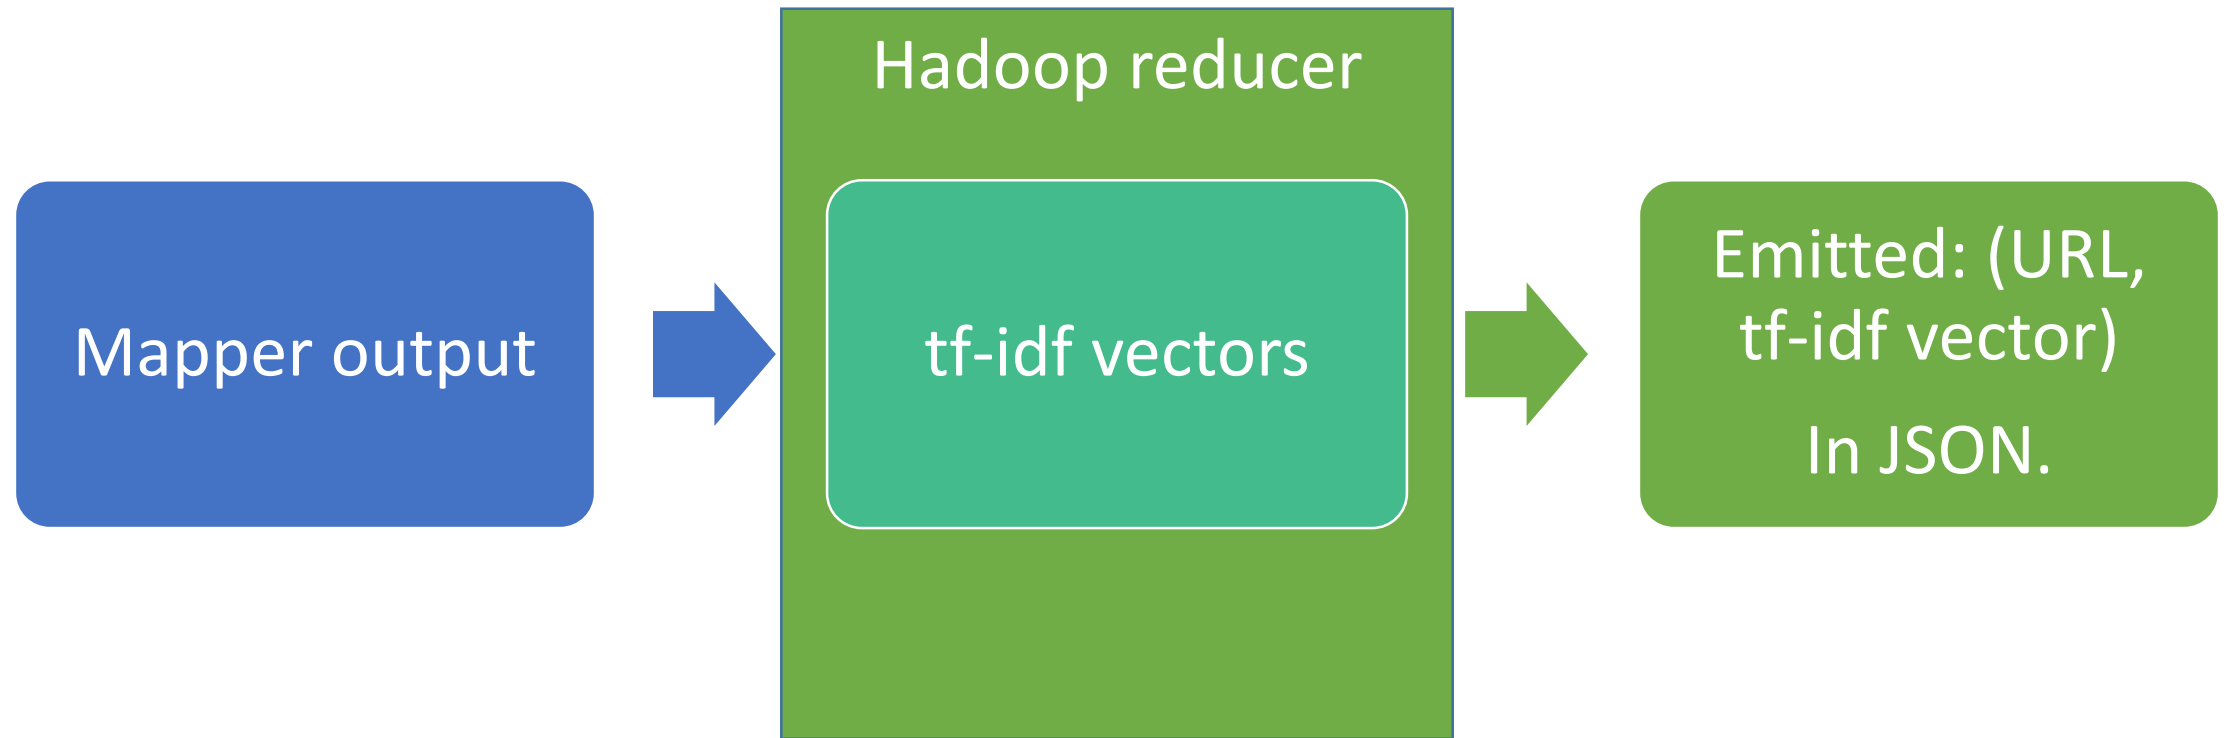
pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



- Document frequency vector is calculated by counting the number of non-zero items on each row.
- From these data structures we can calculate the inverse document frequency (idf) vector.
$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$
- From this, we multiply the normalized frequency matrix with the idf vector to calculate the tf-idf matrix.

How is this parallelized in Hadoop?



Part 3: Clustering and Visualization

Clustering

Visualization

Thank you.

Check out the source on Github: <https://github.com/lnunno/big-data-nyt-tf-idf>

